

Uncertainty-Based Offline Reinforcement Learning with Diversified Q-Ensemble

Advanced Topics in Reinforcement Learning
Winter Semester 2022/23

Outline

1. Introduction - Why Offline RL?
2. Preliminaries
 - a. D4RL: Benchmark for Offline RL
 - b. Actor-Critic Method
 - c. Conservative Q-Learning (CQL) - Baseline
3. SAC-N (Soft Actor-Critic)
4. EDAC (Ensemble-Diversifying Actor-Critic)
5. Comparison

Why Offline RL?




TRIAL & ERROR?

Why Offline RL?

- **Online training** requires active **interaction** with the environment. Interaction with the real world can be **costly and dangerous**
- **Solution:** Offline / “batch” RL - Learn from large, previously collected datasets, **without interaction**
- **Problem:** Out-of-distribution (OOD) actions
 - erroneously high Q-values
 - no feedback from the environment

How to manage uncertainty in offline RL?

Example: OOD actions



- Naive algorithm: treatment causes death?!
- We never see sick patients not treated!

D4RL - Datasets for Deep Data-Driven RL


Datasets specifically designed for benchmarking Offline RL

- Narrow and biased data distribution
- Undirected and multitask data
- Sparse rewards
- Suboptimal data
- Non-representable behavior policies
- Realistic domains


D4RL - MuJoCo Gym environments

- Multi-Joint dynamics with Contact
- physics engine used in robotics, biomechanics, graphics and animation
- standardized in D4RL


Hopper



Walker2D



HalfCheetah






Source: [5]

D4RL - MuJoCo Gym datasets

- expert: train a SAC algorithm online until the strategy reaches the expert performance level, using the expert strategy to collect 1 million samples of data
- medium: first train a SAC algorithm online, stop training in the middle, and then use this partially trained strategy to collect 1 million samples of data
- medium-expert: mix equal amounts of data collected by expert and medium strategies
- medium-replay: train a SAC algorithm online until the strategy reaches a moderate performance level, collecting all the samples placed in the buffer during training
- random: use a random initialization strategy to collect

D4RL - Adroit tasks

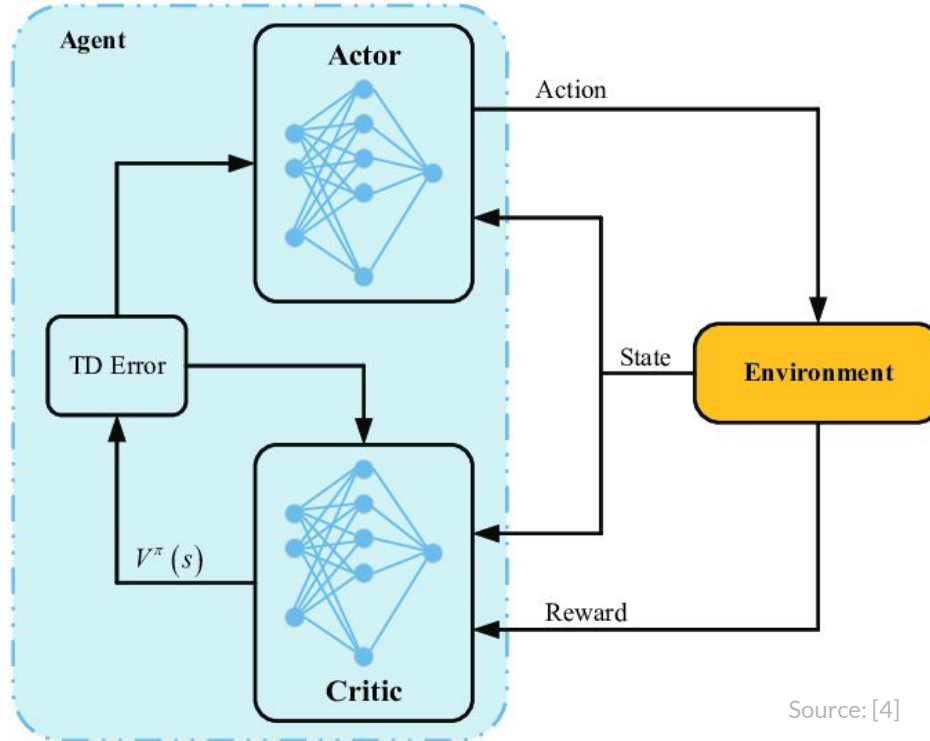
- Realistic robotic manipulation tasks, hand with 24-DoF
- human demonstrations recorded via motion capture
- enables studying human-generated data within a simulated robotic platform
- uses MuJoCo physics simulator



Actor-Critic Method

policy-based actor

value-based critic



Source: [4]

Actor-Critic in Offline RL

Critic Network minimizes

$$J_q(Q_\phi) := \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} \left[\left(Q_\phi(\mathbf{s}, \mathbf{a}) - \left(r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{a}' \sim \pi_\theta(\cdot | \mathbf{s}')} [Q_{\phi'}(\mathbf{s}', \mathbf{a}')] \right) \right)^2 \right]$$

Actor Network maximizes

$$J_p(\pi_\theta) := \mathbb{E}_{\mathbf{s} \sim \mathcal{D}, \mathbf{a} \sim \pi_\theta(\cdot | \mathbf{s})} [Q_\phi(\mathbf{s}, \mathbf{a})]$$

updated in alternating fashion

Conservative Q-Learning (CQL) - Baseline

As of 2021, “Conservative Q-Learning” (CQL) [2] is the state-of-the-art for offline RL. [1]

It uses a “simple Q-value regularizer” to prevent the overestimation of OOD actions.


$$\min_{\phi} J_q(Q_{\phi}) + \alpha \left(\mathbb{E}_{\mathbf{s} \sim \mathcal{D}, \mathbf{a} \sim \mu(\cdot | \mathbf{s})} [Q_{\phi}(\mathbf{s}, \mathbf{a})] - \mathbb{E}_{(\mathbf{s}, \mathbf{a}) \sim \mathcal{D}} [Q_{\phi}(\mathbf{s}, \mathbf{a})] \right)$$

We present two methods **beating the current state-of-the-art**:

- SAC-N: beats CQL
- EDAC: improves runtime performance of SAC-N


SAC-N (Soft Actor-Critic N)

- choose action based on expected reward




SAC-N

- choose action based on expected reward
- penalize uncertainty
 - use the standard deviation as a measure for uncertainty



SAC-N

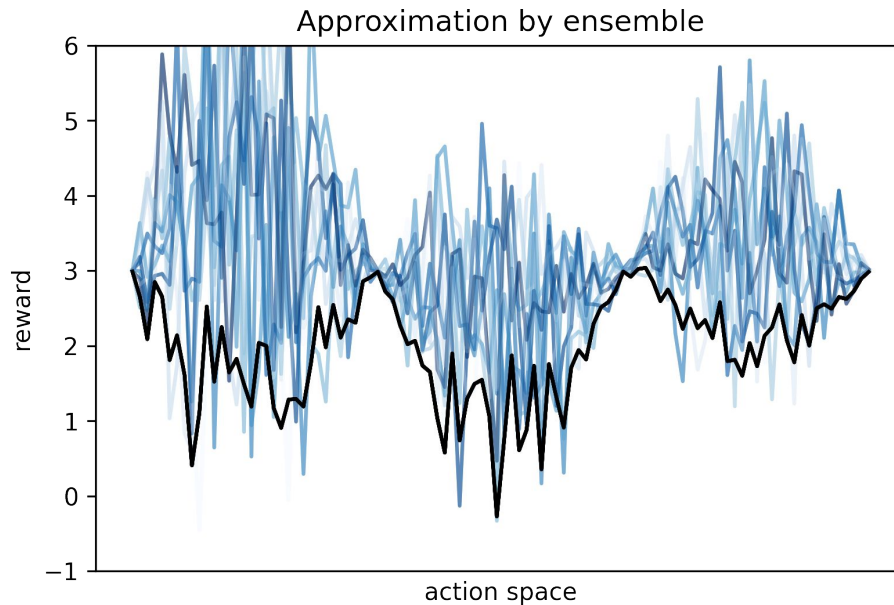
- choose action based on expected reward
- penalize uncertainty
 - **subtract a multiple of the standard deviation** from the expected reward



SAC-N

- choose action based on expected reward
- penalize uncertainty
 - subtract a multiple of the standard deviation from the expected reward
 - **approximate uncertainty by ensemble of Q-networks**

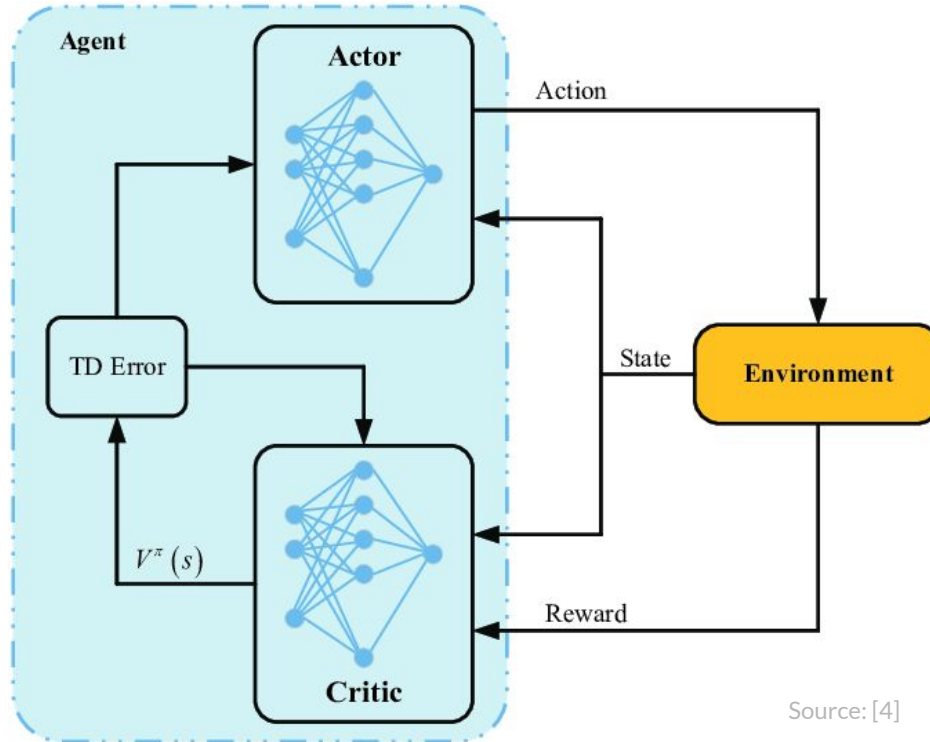
$$\mathbb{E} \left[\min_{j=1, \dots, N} Q_j(\mathbf{s}, \mathbf{a}) \right]$$
$$\approx m(\mathbf{s}, \mathbf{a}) - \Phi^{-1} \left(\frac{N - \frac{\pi}{8}}{N - \frac{\pi}{4} + 1} \right) \sigma(\mathbf{s}, \mathbf{a})$$



Before: Actor Critic Model

policy-based actor

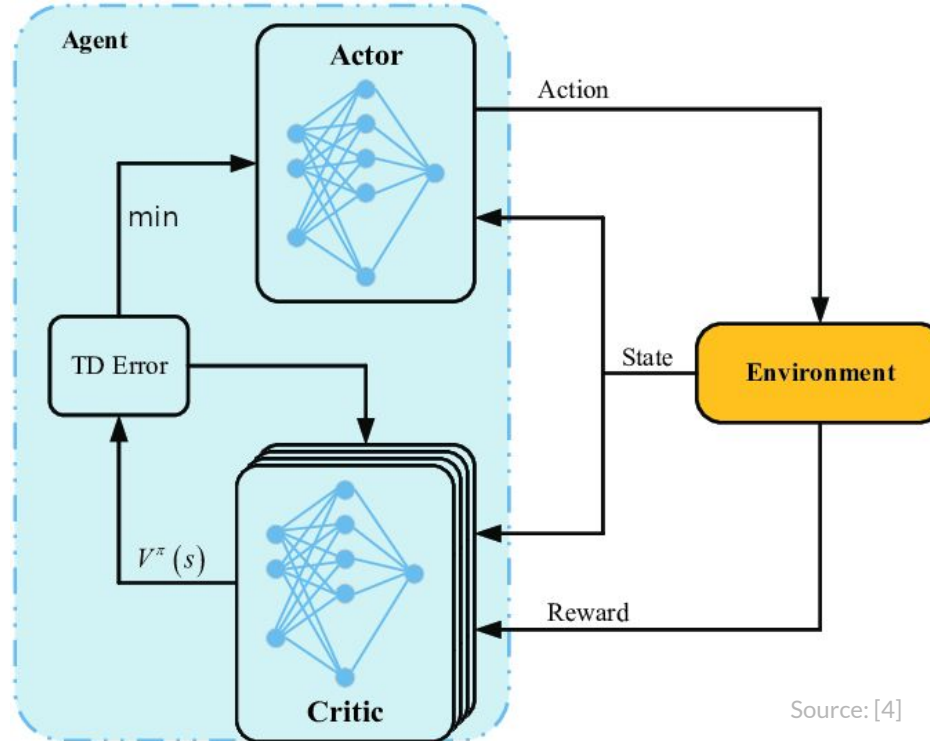
value-based critic



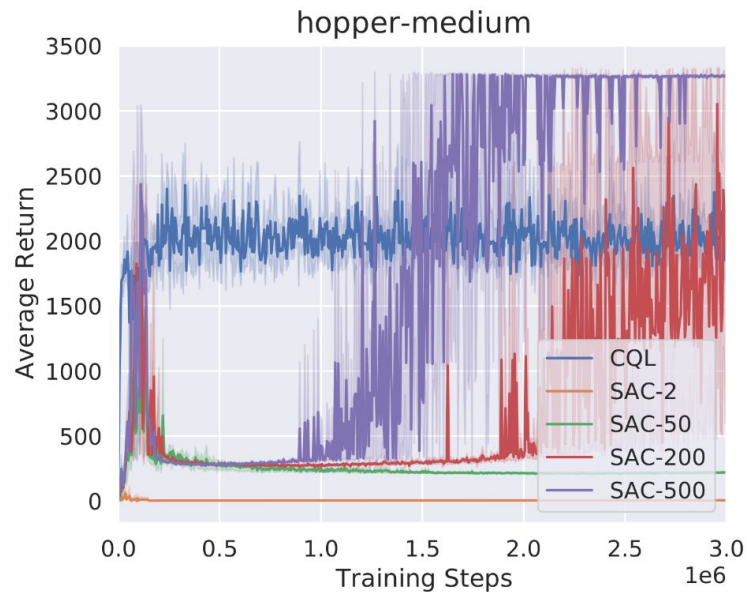

SAC-N

policy-based actor

ensemble of
value-based critics



SAC-N - Results



Source: [1]

SAC-N \rightarrow EDAC


Problem:

A high number of ensemble networks is required, because often many are quite similar to each other.

Idea:


Make sure the ensemble networks are as diverse as possible for OOD actions. Then, fewer networks are required to penalize all possible OOD actions.

EDAC (Ensemble-Diversified Actor Critic)



$\text{Var}(Q_{\phi_j}(s, \mathbf{a} + k\mathbf{w}_2))$ is small so that $\mathbf{a} + k\mathbf{w}_2$ is not sufficiently penalized.

(a) Without ensemble gradient diversification



$\text{Var}(Q_{\phi_j}(s, \mathbf{a} + k\mathbf{w}))$ is large for every direction \mathbf{w} so that all OOD actions are sufficiently penalized.

(b) With ensemble gradient diversification

Source: [1]

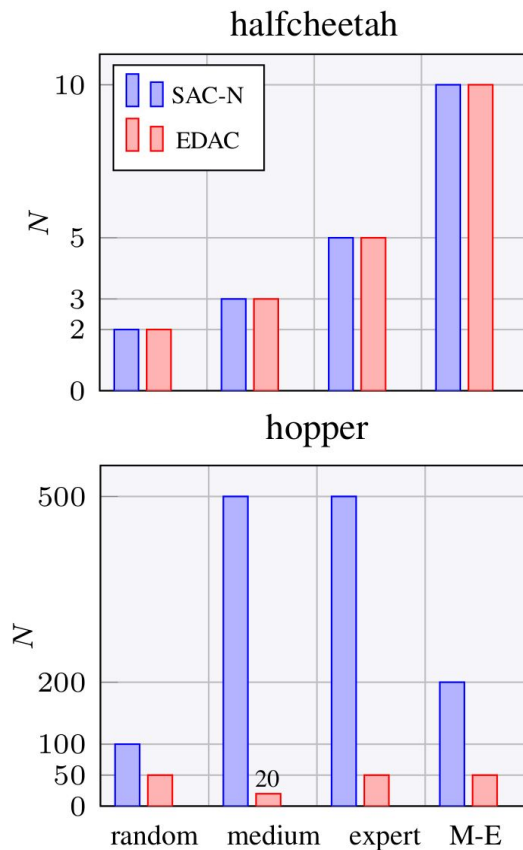
EDAC

Maximize the smallest eigenvalue of the variance of the Q-values for near-distribution OOD actions

$$\underset{\phi}{\text{maximize}} \mathbb{E}_{\mathbf{s}, \mathbf{a} \sim \mathcal{D}} \left[\lambda_{\min} \left(\text{Var} \left(\nabla_{\mathbf{a}} Q_{\phi_j}(\mathbf{s}, \mathbf{a}) \right) \right) \right]$$

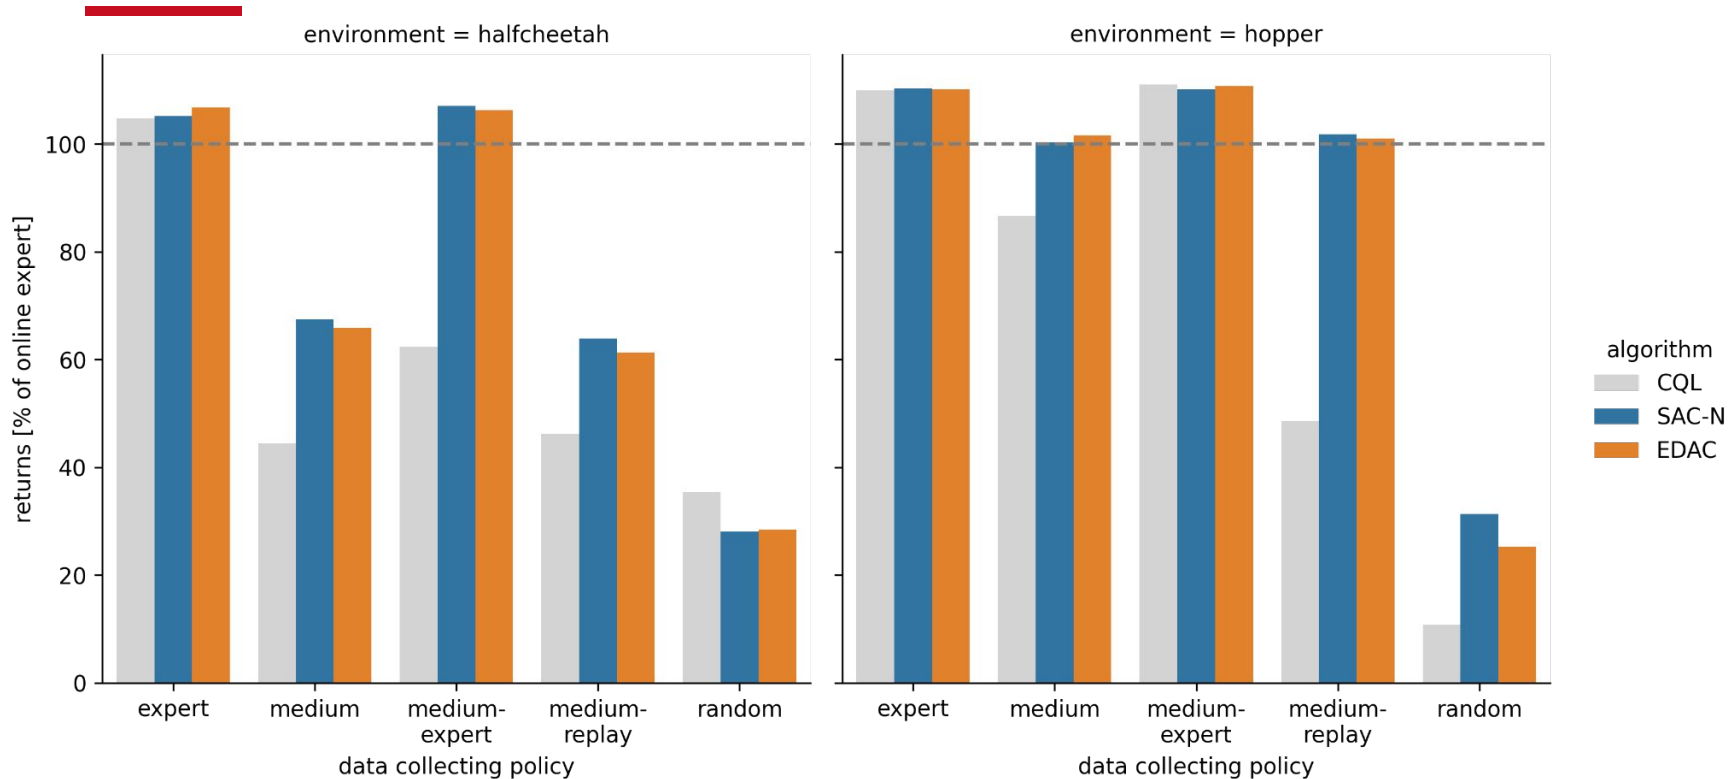
EDAC needs up 90% fewer ensemble networks, compared to SAC-N.

However, this highly varies from task to task.

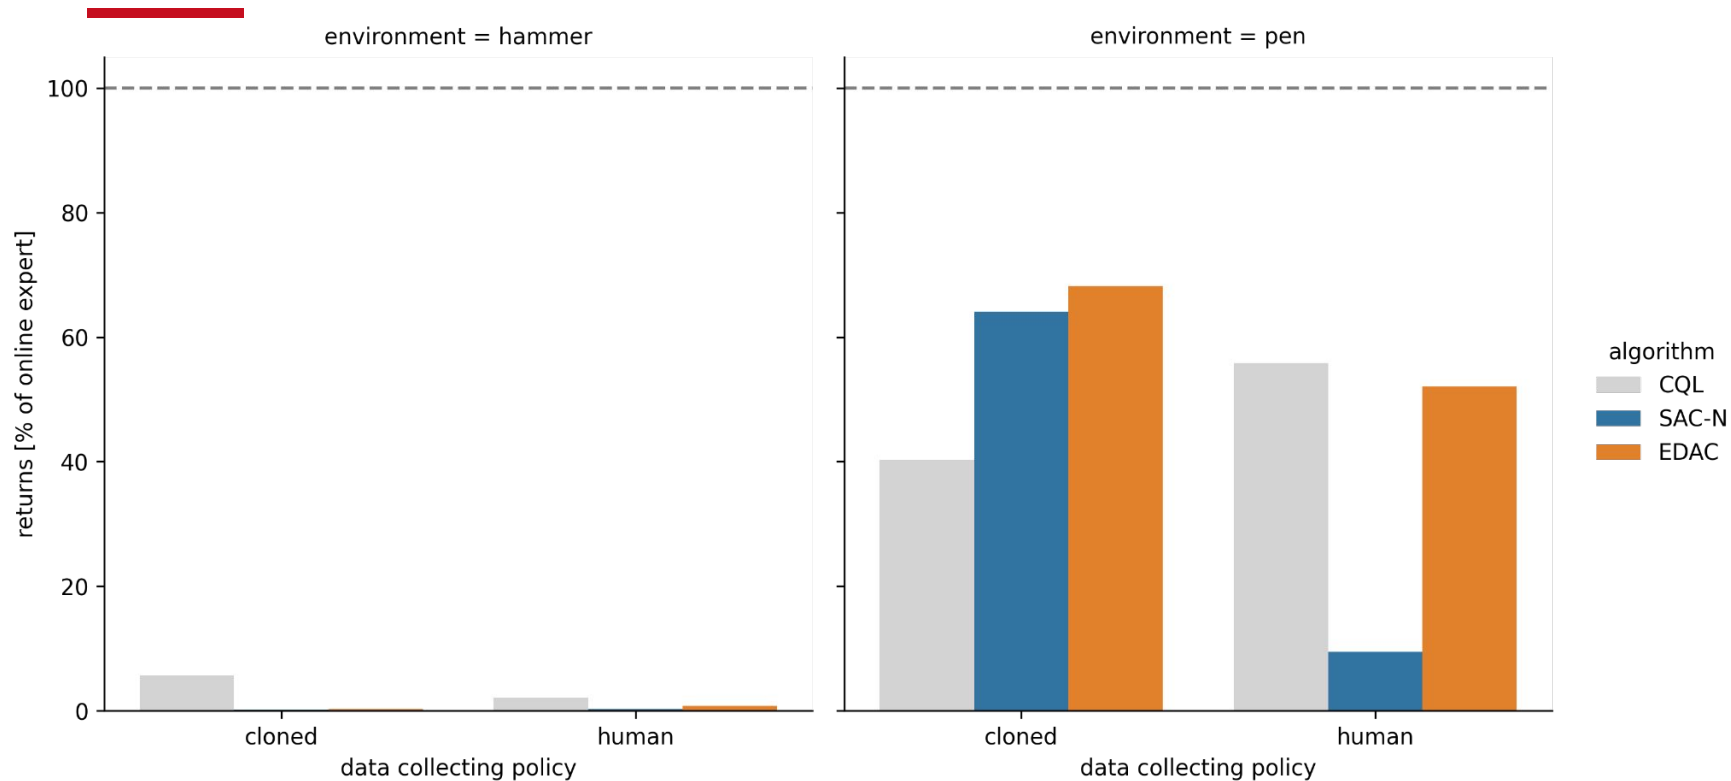


Source: [1]

Comparison: MuJoCo Gym



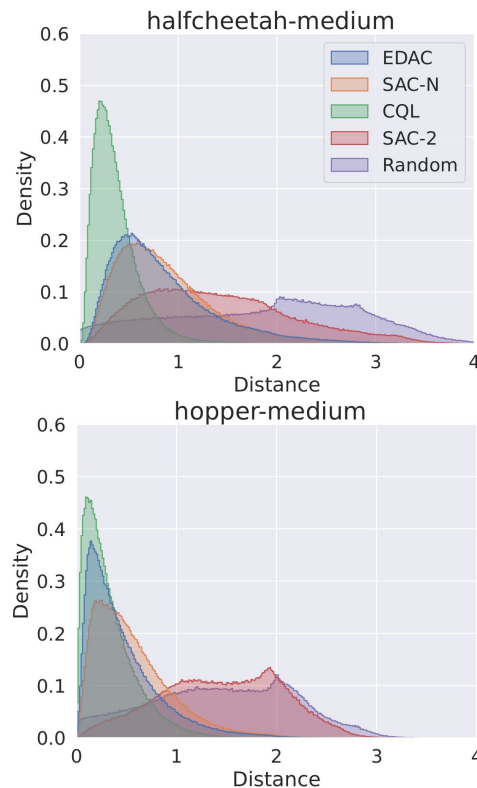
Comparison: Adroit



Comparison: dataset-similarity

SAC-N and EDAC are quite close to the actions in the dataset, but less so than CQL.

Due to the higher average returns, this could indicate a better tradeoff between staying safe, i.e. close to the dataset, and utilizing favorable OOD actions.



Source: [1]

Comparison: Runtime

- EDAC is significantly faster and more memory efficient compared to SAC-500
- EDAC is faster than CQL, but has a higher memory footprint

	Runtime (s/epoch)	GPU Mem. (GB)
SAC	21.4	1.3
CQL	38.2	1.4
SAC-500	44.1	5.1
EDAC	30.8	1.8

Conclusion

- With offline RL, we can avoid trial and error, i.e. dangerous exploration, by using existing data
- A major difficulty are OOD actions, which are not part of the dataset
- We can estimate the uncertainty of the action values und thereby avoid OOD actions by using ensemble networks → SAC-N
- We can reduce the necessary ensemble size by diversification → EDAC
- EDAC outperforms the previous state-of-the-art method (CQL), while being faster (as of 2021)

References

- [1] An, G., Moon, S., Kim, J., & Song, H.O. (2021). Uncertainty-Based Offline Reinforcement Learning with Diversified Q-Ensemble. Neural Information Processing Systems.
- [2] Kumar, A., Zhou, A., Tucker, G., & Levine, S. (2020). Conservative Q-Learning for Offline Reinforcement Learning. ArXiv, abs/2006.04779.
- [3]
- [4] Giang, Hoang & Hoan, Tran & Thanh, Pham & Koo, Insoo. (2020). Hybrid NOMA/OMA-Based Dynamic Power Allocation Scheme Using Deep Reinforcement Learning in 5G Networks. Applied Sciences. 10. 4236. 10.3390/app10124236.
- [5] Justin Fu (2020). D4RL: Building Better Benchmarks for Offline Reinforcement Learning. <https://bair.berkeley.edu/blog/2020/06/25/D4RL/>