




Verifiable Reinforcement Learning via Policy Extraction

Osbert Bastani, Yewen Pu, Armando Solar-Lezama




Verifiable Reinforcement Learning via Policy Extraction

Claim: Verify that our RL agent is safe!




Katz et al. (2017)



DNN agent:

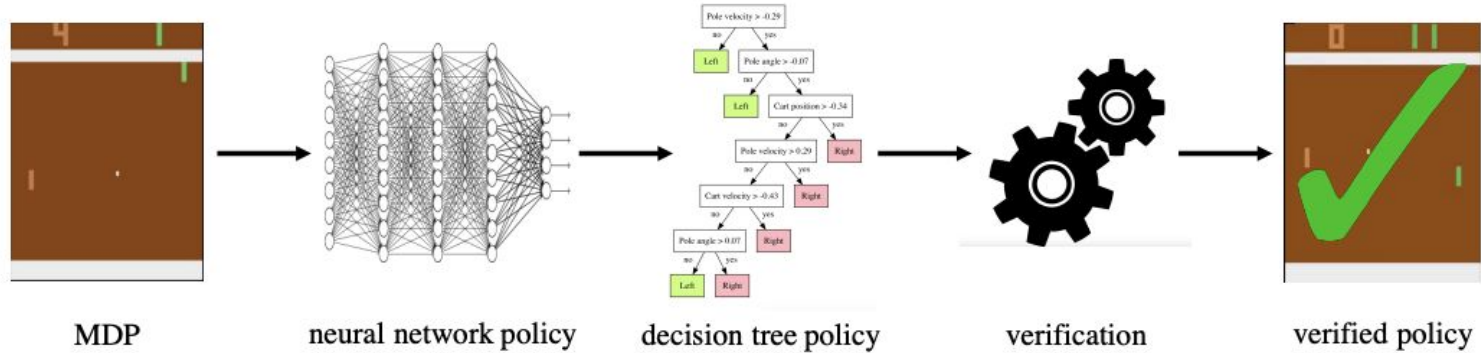
- Easy to train
- Hard to verify



Tree agent:

- Hard to train
- Easy to verify

How to get a verifiable RL policy?




Bastani et al. (2019)

Structure of The Talk


1. Policy Extraction
2. Policy Verification
 - a. Correctness
 - b. Robustness
 - c. Stability
3. Evaluation
4. Discussion

Policy Extraction via Imitation Learning




Ross and Bagnell (2011)

Policy Extraction via DAgger




Ross and Bagnell (2011)

Idea: The imitator should focus on critical states



$$V_t^{(\pi^*)}(s) \approx \min_a Q_t^{(\pi^*)}(s, a)$$



$$V_t^{(\pi^*)}(s) \gg \min_a Q_t^{(\pi^*)}(s, a)$$



VIPER: Sampling

Verifiability via Iterative Policy ExtRaction

Define this measure of “criticalness” of a state

$$\tilde{\ell}_t(s) = V_t^{(\pi^*)}(s) - \min_{a \in A} Q_t^{(\pi^*)}(s, a)$$

And use it to re-sample from our trace data:

$$(s, a) \sim p((s, a)) \propto \tilde{\ell}_t \mathbb{I}[(s, a) \in D]$$



VIPER: Algorithm

Verifiability via Iterative Policy ExtRaction

Algorithm 1 Decision tree policy extraction.

```
procedure VIPER( $(S, A, P, R), \pi^*, Q^*, M, N$ )  
  Initialize dataset  $\mathcal{D} \leftarrow \emptyset$   
  Initialize policy  $\hat{\pi}_0 \leftarrow \pi^*$   
  for  $i = 1$  to  $N$  do  
    Sample  $M$  trajectories  $\mathcal{D}_i \leftarrow \{(s, \pi^*(s)) \sim d^{(\hat{\pi}_{i-1})}\}$   
    Aggregate dataset  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$   
    Resample dataset  $\mathcal{D}' \leftarrow \{(s, a) \sim p((s, a)) \propto \tilde{\ell}(s) \mathbb{I}[(s, a) \in \mathcal{D}]\}$   
    Train decision tree  $\hat{\pi}_i \leftarrow \text{TrainDecisionTree}(\mathcal{D}')$   
  end for  
  return Best policy  $\hat{\pi} \in \{\hat{\pi}_1, \dots, \hat{\pi}_N\}$  on cross validation  
end procedure
```



VIPER: Theoretical guarantees

Theorem 2.2. *For any $\delta > 0$, there exists a policy $\hat{\pi} \in \{\hat{\pi}_1, \dots, \hat{\pi}_N\}$ such that*

$$J(\hat{\pi}) \leq J(\pi^*) + T\varepsilon_N + \tilde{O}(1)$$

with probability at least $1 - \delta$, as long as $N = \tilde{\Theta}(\ell_{\max}^2 T^2 \log(1/\delta))$.

$$\tilde{\ell}_t(s, \pi) = \tilde{\ell}_t(s) \tilde{g}(s, \pi)$$

Implies that we can achieve the same training loss via re-sampling:

$$\mathbb{E}_{(s,a) \sim p((s,a))} [\tilde{g}(s, \pi)] = \mathbb{E}_{(s,a) \sim \mathcal{D}} [\tilde{\ell}(s, \pi)]$$

1. Policy Extraction
- 2. Verifying the Decision Tree Policy**
 - a. Correctness
 - b. Robustness
 - c. Stability
3. Evaluation
4. Discussion

Correctness for Toy Pong

$$f_{\pi}(s) = f_i(s) = \beta_i^T s$$

$$\psi = \left(\bigwedge_{t=1}^{t_{\max}} \phi_t \right) \wedge \psi_0 \Rightarrow \bigvee_{t=1}^{t_{\max}} \psi_t$$


$$\psi_t = (s_t \in Y_0)$$

ϕ_t : Inductive controller invariant

Controller is correct when $\neg\psi$ cannot be satisfied


Correctness for Toy Pong

- 30 Decision tree nodes vs 700 NN neurons
- SMT solved in <3 seconds
- Finds policy error!



Bastani et al (2019)

Robustness for Toy Pong



VIPER:

- Completes in seconds
- Accurate to ϵ within 10-5

Reluplex:

- Huge variance of completion times
- One timeout even
- Accurate to ϵ within 0.1

Stability for cartpole

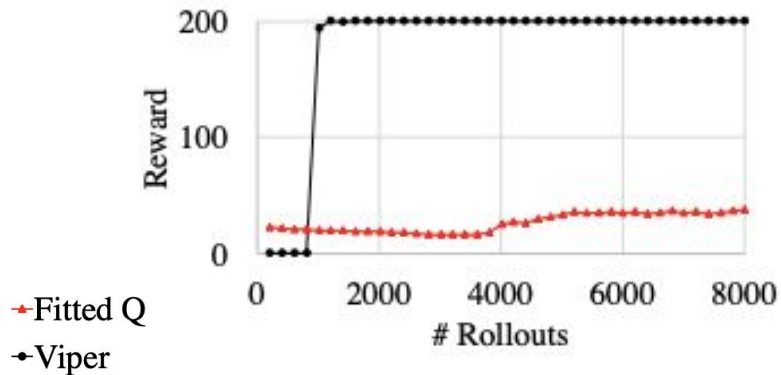
- Uses an iLQR oracle
- Achieves perfect reward on Cartpole
- Three node tree with linear regressors

Evaluation:

- VIPER is verified at stability region with L_{∞} norm ≤ 0.03 in 4 seconds
- NN requires enumeration which takes 10 min. and verifies area 10^{-15} of stability region


Comparing VIPER to other methods

Cartpole



Vs FittedQ

Toy Pong



DAgger

Discussion

- Policy Extraction also useful for explainable AI.
- Extracted policies need manual fixing.
- Verification process requires many approximations.
- What is the limit to decision tree extraction?

Things to take away if nothing else

- You can efficiently distill a trained DNN agent into a decision tree and have theoretical upper bounds on its training reward.
- The key idea of VIPER is sampling the Oracle in such a way that critical states are given more important weight ($Q_{\text{opt}} \gg Q_{\text{worst}}$).
- Using a decision tree policy you can efficiently verify attributes such as correctness, stability, robustness.

Sources

- Bastani et al (2019)
- Ross et Bagnell (2011)
- <https://trustml.github.io/docs/viper-presentation.pdf>
- Katz et al. (2017)

Backup slides

How do you obtain the loss for continuous actions, i.e. when you cannot find Q_{\min} ?

Instead, we used an approach inspired by guided policy search [21]. We trained another decision tree using a different oracle, namely, an iterative linear quadratic regulator (iLQR), which comes with stability guarantees (at least with respect to the linear approximation of the dynamics, which are a very good near the origin). Note that we require a model to use an iLQR oracle, but we anyway need the true model to verify stability. We use iLQR with a time horizon of $T = 50$ steps and $n = 3$ iterations. To extract a policy, we use $Q(s, a) = -J_T(s)$, where $J_T(s) = s^T P_T s$ is the cost-to-go for the final iLQR step. Because iLQR can be slow, we compute the LQR controller for the linear approximation of the dynamics around the origin, and use it when $\|s\|_\infty \leq 0.05$. We now use continuous actions $A = [-a_{\max}, a_{\max}]$, so we extract a (3 node) decision tree policy π with linear regressors at the leaves (internal branches are axis-aligned); π achieves a reward of 200.0.

Problem formulation

Problem formulation. Let (S, A, P, R) be a finite-horizon (T -step) MDP with states S , actions A , transition probabilities $P : S \times A \times S \rightarrow [0, 1]$ (i.e., $P(s, a, s') = p(s' \mid s, a)$), and rewards $R : S \rightarrow \mathbb{R}$. Given a policy $\pi : S \rightarrow A$, for $t \in \{0, \dots, T-1\}$, let

$$V_t^{(\pi)}(s) = R(s) + \sum_{s' \in S} P(s, \pi(s), s') V_{t+1}^{(\pi)}(s')$$

$$Q_t^{(\pi)}(s, a) = R(s) + \sum_{s' \in S} P(s, a, s') V_{t+1}^{(\pi)}(s')$$

be its value function and Q -function for $t \in \{0, \dots, T-1\}$, where $V_T^{(\pi)}(s) = 0$. Without loss of generality, we assume that there is a single initial state $s_0 \in S$. Then, let

$$\begin{aligned} d_0^{(\pi)}(s) &= \mathbb{I}[s = s_0] \\ d_t^{(\pi)}(s) &= \sum_{s' \in S} P(s', \pi(s'), s) d_{t-1}^{(\pi)}(s') \quad (\text{for } t > 0) \end{aligned}$$

be the distribution over states at time t , where \mathbb{I} is the indicator function, and let $d^{(\pi)}(s) = T^{-1} \sum_{t=0}^{T-1} d_t^{(\pi)}(s)$. Let $J(\pi) = -V_0^{(\pi)}(s_0)$ be the cost-to-go of π from s_0 . Our goal is to learn the best policy in a given class Π , leveraging an *oracle* $\pi^* : S \rightarrow A$ and its Q -function $Q_t^{(\pi^*)}(s, a)$.

Reward bound

Theorem 2.2. *For any $\delta > 0$, there exists a policy $\hat{\pi} \in \{\hat{\pi}_1, \dots, \hat{\pi}_N\}$ such that*

$$J(\hat{\pi}) \leq J(\pi^*) + T\varepsilon_N + \tilde{O}(1)$$

with probability at least $1 - \delta$, as long as $N = \tilde{\Theta}(\ell_{\max}^2 T^2 \log(1/\delta))$.

In contrast, the bound $J(\hat{\pi}) \leq J(\pi^*) + uT\varepsilon_N + \tilde{O}(1)$ in [25] includes the value u that upper bounds $Q_t^{(\pi^*)}(s, a) - Q_t^{(\pi^*)}(s, \pi^*(s))$ for all $a \in A$, $s \in S$, and $t \in \{0, \dots, T - 1\}$. In general, u may be $O(T)$, e.g., if there are *critical states* s such that failing to take the action $\pi^*(s)$ in s results in forfeiting all subsequent rewards. For example, in cart-pole [5], we may consider the system to have failed if the pole hit the ground; in this case, all future reward is forfeited, so $u = O(T)$.

An analog of u appears implicitly in ε_N , since our loss $\tilde{\ell}_t(s, \pi)$ includes an extra multiplicative factor $\tilde{\ell}_t(s) = V_t^{(\pi^*)}(s) - \min_{a \in A} Q_t^{(\pi^*)}(s, a)$. However, our bound is $O(T)$ as long as $\hat{\pi}$ achieves high accuracy on critical states, whereas the bound in [25] is $O(T^2)$ regardless of how well $\hat{\pi}$ performs.

Correctness for Toy Pong

We partition the state so that for every partition S_i we can get a β_i from the policy.

$$f_\pi(s) = f_i(s) = \beta_i^T s$$

Either we were not in this state previously or our current state is a result from the dynamics

$$\phi_t = \bigvee_{i=1}^k (s_{t-1} \in S_i \Rightarrow s_t = \beta_i^T s_{t-1}) \quad \forall t \in \{1, \dots, t_{\max}\}$$

$$\psi = \left(\bigwedge_{t=1}^{t_{\max}} \phi_t \right) \wedge \psi_0 \Rightarrow \bigvee_{t=1}^{t_{\max}} \psi_t$$